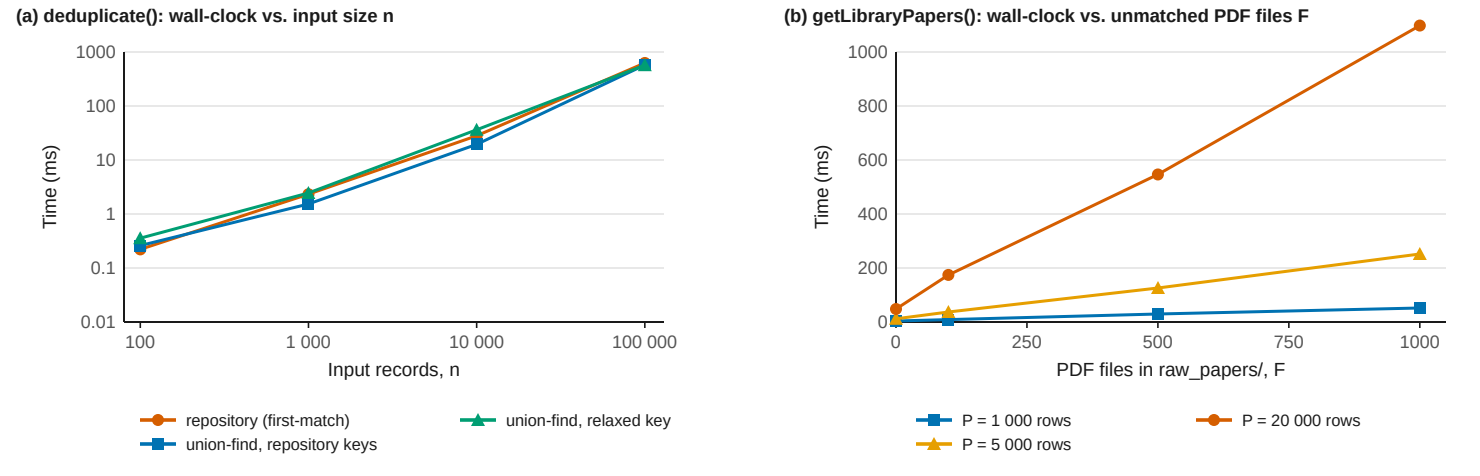
(a) deduplicate(): wall-clock vs. input size n
Time (ms)
1000
100
10
1
0.1
0.01
100
1 000
10 000
100 000
Input records, n
repository (first-match)
union-find, repository keys
union-find, relaxed key
(b) getLibraryPapers(): wall-clock vs. unmatched PDF files F
Time (ms)
1000
800
600
400
200
0
0
250
500
750
1000
PDF files in raw_papers/, F
P = 1 000 rows
P = 5 000 rows
P = 20 000 rows